



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Lecture with Computer Exercises:
Modelling and Simulating Social Systems with MATLAB

Project Report

**Evacuation Bottleneck
Simulating a Panic on a Cruise Ship**

Benedek Vartok & Johannes Weinbuch

Zurich
December 2009

Agreement for free-download

We hereby agree to make our source code for this project freely available for download from the web pages of the SOMS chair. Furthermore, we assure that all source code is written by ourselves and is not violating any copyright restrictions.

Johannes Weinbuch

Benedek Vartok



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Declaration of Originality

This sheet must be signed and enclosed with every piece of written work submitted at ETH.

I hereby declare that the written work I have submitted entitled

Evacuation Bottleneck — Simulating a Panic on a Cruise Ship

is original work which I alone have authored and which is written in my own words.*

Author(s)

Last name

Meinbach

Vartok

First name

Johannes

Benedek

Supervising lecturer

Last name

Balietti

Donnay

First name

Stefano

Karsten

With the signature I declare that I have been informed regarding normal academic citation rules and that I have read and understood the information on 'Citation etiquette' (http://www.ethz.ch/students/exams/plagiarism_s_en.pdf). The citation conventions usual to the discipline in question have been respected.

The above written work may be tested electronically for plagiarism.

Zürich, 20.12.12-10
Place and date

J. Meinbach Benedek Vartok
Signature

*Co-authored work: The signatures of all authors are required. Each signature attests to the originality of the entire piece of written work in its final form.

Print form

Contents

1	Abstract	5
2	Individual Contributions	5
3	Introduction and Motivations	5
4	Description of the Model	5
5	Implementation	6
5.1	Input	6
5.2	The Simulation Routines	7
5.2.1	Code Reuse from Multilevel Evacuation	7
5.2.2	General Structure	8
5.2.3	Main Loop	8
5.2.4	Agent Placement	8
5.2.5	Agent Dynamics	8
5.2.6	Removing Filled Exits	10
5.3	Output and Plotting	11
6	Simulation Results and Discussion	11
6.1	Passenger distribution	11
6.2	Panic Level	12
6.3	Closed Exits	14
7	Summary and Outlook	15
8	References	15

1 Abstract

This work takes a look into the evacuation mechanisms of a cruise ship in case of an emergency. A simple model is implemented which is used to simulate the dynamics of such a system. The main emphasis was on the limited capacity of the exits, since that is the key element for a rescue boat.

2 Individual Contributions

The work on this project was split among us to fit our strengths the best way possible. Because of his knowledge in image editing and formats, Johannes Weinbuch focused on the image manipulation for the input and implemented the loading of the image into MATLAB, improving the existing solutions from the previous courses. He further took a large part of the writing for the report and executing the simulations, which were written by Benedek Vartok. He evaluated which code from previous semesters could and should be reused, and implemented the missing parts for our special case. Also, he wrote the output mechanisms for the simulation, so that the data could be used for analysis.

3 Introduction and Motivations

In January 2012, the Costa Concordia hit a rock and ran aground[1]. This event got great media attention for a long time so we decided to take a closer look at the evacuation of a cruise ship. The question is, what is the best strategy to leave the ship? This question should for sure be answered with one of the emergency drills, but it is always good to have some background knowledge.

So, our key questions are: Which is the best strategy for evacuation concerning the choice of the way towards the rescue boats. Should all passengers distribute equally over the entries, or is there a better one? Also, which one takes longer: a high panic level on a nearly empty ship or a low panic level on a very full ship. What happens, if a boat suddenly is inoperable? How can the reaction be optimized?

4 Description of the Model

The model is a big simplification of real life, otherwise it would be way too complex to simulate. It assumes that the ship is intact, that there is calm sea and that the passengers are obliged to leave the ship. A possible explanation for this could be a machine defect which leaks explosive gas in a badly ventilated room in the ship.

Further, we assume that the rescue boats are like doors, which close after a certain amount of people going through them.

Since we also assume that the other doors, for example between the rooms or floors, are constantly open and working, we only simulate one deck, the one with the exits to the rescue boats. The evacuation of multiple floors in a static building has already been researched in [2].

After these simplifications, the task left to simulate was the evacuation of a single floor with some elements that can change. For this task, we chose a simple agent based modeling solution as described in [3]. A passenger is treated as a particle. It has a mass, and there are physical and social forces, accelerating that mass so that it cannot always follow its desired direction. The desired direction is implemented as the shortest path to the nearest exit. For the exact formulae for the forces see section 5.2.5.

The floor is given by a deck plan for the Costa Voyager [4]. Since these deck plans usually are made for advertisement purposes, it is to be expected that they are not absolutely accurate. So during image processing for simplification of this plan, a few further assumptions were made, mainly about the actual sizes and capacities. We won't list them all here, because they are also in the configuration files.

The last element of additional complexity was added to get an idea of how the closing of an exit works. A switch was added to decide whether every agent should know instantaneously about the closed exit or if the knowledge spreads over time. The technical aspect of this is further explained in section 5.2.6.

5 Implementation

5.1 Input

Since we had some good projects which covered similar problems as ours, we could get some ideas from them, but at the same time improve them. Namely, there are [2] and [5]. As far as the input for the simulation is concerned, we see two approaches in these works for getting the map data into the simulation. In [2], a simple PNG image is used to get a map into the simulation. The problem here is that only a certain RGB color value can be read out of the image. This can lead to problems if the image is processed with automatic or semiautomatic image manipulation programs, since only a minor difference in color can prevent the generation of the desired data. In [5], the image format is even more simple. There is only a bitmap image read into MATLAB. Since the bitmap images can use a colormap, MATLAB doesn't use 3 channels but a unique number for each color in an image matrix to give every pixel its color. This has the same problem as the PNG solution regarding how exactly the

colors have to be set, but the different parts of the image can be separated with less code.

We took the best of both solutions. We used the PNG-format with indexed colors. So we have the most flexibility with very little usage of disk space. There is no special “wall color” or anything like that, just a simple rule how the colormap is read: Color 0 of the map specifies walls, color 1 free space. Then, there can be any number of spawn zones. Spawn zones are the areas in the image, where new agents can be placed. With different spawn zones, it is possible to account for different situations: A ballroom is different from a staircase. The number of spawn zones is specified in the configuration file. At last, there is an arbitrary number of exits. Again, each exit can have its own parameters or can even be handled specially in the program’s code.




Figure 1: Screenshot of the Rearrange Colormap dialog in Gimp 2.6.11

To manipulate the colormap, any slightly sophisticated image manipulation program should suffice. We used the free software Gimp [6]. It has a very convenient command which allows the user to rearrange the colormap. This is shown in figure 1.

5.2 The Simulation Routines

5.2.1 Code Reuse from Multilevel Evacuation

Since this project has very similar foundations as [2] (like the forces used in the model), we were able to use a lot of code and some design decisions from their program. Some of the structure needed to be changed to implement our custom features, but for example the utility functions for the Fast Sweeping algorithm (`fastSweeping.c`), generating gradients (`getNormalizedGradient.c`) and the linear interpolation (`lerp2.c`) which the other group wrote in C were copied without modification into our code-tree.

5.2.2 General Structure

During the entire program run, one single big structure is used to hold and pass around the state of the simulation.

At first, this structure is initialized with the fields that are given in the configuration file by `loadConfig`¹. Then, `initialize` runs over the struct and calculates some data needed in the simulation, such as the vector fields needed for the wall and exit force fields using the `fastSweeping` method which we got from [2].

5.2.3 Main Loop

`simulate` is the routine we call when we do an entire simulation. It uses the `loadConfig` and `initialize` functions to initialize its runtime data, then it does the loop calculating forces, adding new agents, progressing the agents, updating the exit vector fields, potentially plotting and saving frames and collecting data.

These steps will be explained in detail in the following sections.

5.2.4 Agent Placement

As mentioned in section 5.1, our model of the ship has different spawning zones where new agents can start out at. The way we implemented it, in every step of the simulation loop it is checked whether there are any remaining agents that need to be placed (i.e. agents which are not in the simulation yet). If there are, then for every agent the program chooses a random point in the spawning zones and places him there, unless it detects that the agent would collide either with walls or other agents. In that case, the routine tries the placement for that agent up to five times, each time with a new random position. If the agent couldn't be placed, then he will have a chance to spawn in the next time step.

This method was implemented in `placeAgents`. In the same method, the basic properties of the agent get assigned, like the starting zero velocity and a random radius.

5.2.5 Agent Dynamics

To simulate the movement of the agents in this physical model, different forces need to be calculated in every step on every agent. All of the force formulae have been taken from [3]. These forces have been separated into the following functions which are called by `simulate`:

¹The MATLAB implementation of our functions is in the `code` directory: `loadConfig` can be found in the file `code/loadConfig.m` etc.

addDesiredForces is responsible for making the agents seek the exits of the layout, in our case the rescue boats.

This is accomplished by giving an agent a “desired” velocity vector pointing along the shortest path to the nearest exit. The vector is sampled and interpolated (using `lerp2` from [2]) from a vector field which is calculated at the beginning (see 5.2.2) of the simulation and when the exits change (see 5.2.6).

Using this desired vector \vec{e} we can say what force addition the agent gets:

$$\vec{F}_{\text{desired}} = m \frac{v_0 \vec{e} - \vec{v}}{\tau}$$

where m is the agent’s mass, v_0 is his target speed, \vec{v} is his current velocity vector and τ is a characteristic time determining how fast the desired velocity should be reached.

addInterAgentForces models the repulsive forces between agents: They do not want to get too close together and if they touch, they have to be kept apart physically and some friction appears. The model we use has this formula:

$$\vec{F}_{\text{agents}} = (Ae^{(r-d)/B} + k \max\{0, r - d\})\vec{n} + \kappa \max\{0, r - d\}\Delta\vec{v} \cdot \vec{t}$$

where r is the sum of radii of both agents, d is the distance between the center points of the agents, \vec{n} is the normalized vector between the two agents, \vec{t} is the tangential vector and $\Delta\vec{v}$ is the velocity difference vector. A influences the magnitude of the “social” repulsive force, B specifies a factor for the range of influence for this force, k gives the strength of the physical separation force and κ is a friction coefficient.

For finding possible agent pairs to calculate the function on we used the naive approach of checking every possible pairing with two nested loops and then using a cutoff distance to avoid calculating this complicated force expression when it’s too small to matter anyways. This method has complexity $O(N_{\text{agents}}^2)$ which is far from optimal; we also tested the Range Tree implementation of [2] for our program, however benchmarks didn’t show a noticeable gain in efficiency.

addWallForces calculates agents avoiding and experiencing resistance from walls. Just like agent-agent repulsion, this force has a “social”, a physical and a frictional component:

$$\vec{F}_{\text{walls}} = (Ae^{(r-d)/B} + k \max\{0, r - d\})\vec{n} - \kappa \max\{0, r - d\}(\vec{v} \cdot \vec{t})\vec{t}$$

where r is the radius of the agent, d is his distance from the wall, \vec{n} is the wall normal vector, \vec{t} is the wall tangent vector and \vec{v} is the agent's velocity vector. The coefficients are the same as in F_{agents} .

Accessing the wall distances and normals is done similarly to `addDesiredForces`, with precalculated fields using the Fast Sweeping method.

`progressAgents` applies the forces from the above listed functions, using them to update the agents' positions and velocities for the next simulation step.

To accomplish this, the leap-frog integration scheme was used. In every step, the following recalculations of the velocities and positions of the agents take place:

$$\begin{aligned}\vec{v} &= \vec{v} + \Delta t \cdot \frac{\vec{F}}{m} \\ \vec{x} &= \vec{x} + \Delta t \cdot \vec{v}\end{aligned}$$

This method is fairly well suited for physical simulations with forces such as these. However, extra measures were taken to improve the stability of our program: Before doing any further calculations with them, the velocities and forces get clipped to a configurable maximum magnitude to avoid instabilities for too high step-sizes or too strong force parameters.

Without this, the simulation of some agents might get out of control if they get too close to walls or to other agents and behave in unphysical ways.

This function also checks whether some agents have entered a non-full exit zone (a rescue boat) and if the exit got full from them, it closes that one. What exactly happens then is explained in the next section.

5.2.6 Removing Filled Exits

In our model we have several exit zones, each with a maximum capacity (since they are rescue boats with limited size). Because of that the simulation needs to take it into account when one of them gets filled: It needs to remove the exit and update the vector fields for the desired velocities which the agents use to find the nearest exit. For this update we modeled and implemented two different approaches which can be chosen in the configuration file.

The first one simply recalculates the entire vector field when an exit is closed with the Fast Sweeping method, just like in the initialization, but with the full exits removed. The update is instantaneous, so every agent on the entire ship reacts to

it immediately. This is somewhat unrealistic, which is why we came up with the second method.

Instead of updating the entire vector field right away, we can just calculate the updated field and use that to update growing regions. To be exact, when an exit closes, a circle starts growing around at a configurable rate, and in this area the new destination vector field is used. That way, at first only agents close to that exit react to the change, then over time the more distant agents also “notice” the filled exit and go for a different one.

The circle-shaped destination field update is implemented in `progressDestFields`.

5.3 Output and Plotting

Our program has two plotting functions:

`plotFloor` draws the ship’s layout and all the agents on it in the current simulation state. It is called from `simulate` in every loop iteration and the resulting picture is saved to `code/frames/`, but only if the option for saving frames has been enabled in the configuration file.

`plotExitedAgents` is called at the end of the simulation and creates time series plots of the rescue boat occupations.

Also the program saves its entire state data object to a file in `code/frames/` at the end, which includes other information as well, such as the time needed for all agents to reach the exits and the time series of the total escaped agents.

6 Simulation Results and Discussion

6.1 Passenger distribution

If we plot the exited agents over a time axis for each exit, we can see that the boats tend to be filled in a sequential way (See Figure 2 and 3). This means, a new boat mostly gets frequented if the old boat is already full. This effect is very strong and good to see if we have only few passengers, but also if there are many, a tendency towards this can be observed. The only thing, that interferes with this observation is, that at the beginning, two boats are filled at the same time. This is because some agents are placed so that they have a shorter distance to the left.

Although there could be numerous reasons for this to happen, in reality a more distributed and parallel filling of the boats is expected. The simulation is based on a model which takes the distance to the nearest exit as the indicator for the desired direction. In reality, people want to get out of the ship the fastest way possible, especially if there are only limited capacities on the boats. This suggests that the nearest exit is not the best strategy to get out fast. On the specific geometry of




Figure 2: Plot of the filling of the rescue Boats on a ship with 700 Agents, $v_0 = 1.5 \frac{m}{s}$. Each line represents a different rescue boat.

the Costa Concordia, the corridor from bow to stern ends more on the starboard (right) side. The agents follow the shortest path and start to jam up, while on the left there are no obstacles. Without regard to the beginning, two boats get used in parallel only if there are many agents. This can be explained by the jam, so people get pushed back towards the other boat. Also, if a boat is empty, the Crowd can separate into two parts: One, that is nearer to the next boat on the same side and one, which was further away from the empty boat, which now has the nearest boat on the other side.

The expectation for reality is that people would recognize that the nearest boats take more time to reach than the ones further away. Thus, they would distribute more equally over the different boats.

6.2 Panic Level

As described in [3], a bigger desired velocity v_0 can lead to longer times in evacuation. At a higher panic level, people want to go faster, but it can take longer to get out. When v_0 is below 1.5 m/s , we could reproduce the results of decreasing evacuation times both with many and few agents, as seen on Figures 4 and 5.

However, with higher values for v_0 , we could not reproduce the results, because agents got pushed into walls and remained stuck there. In Figure 4, we had two

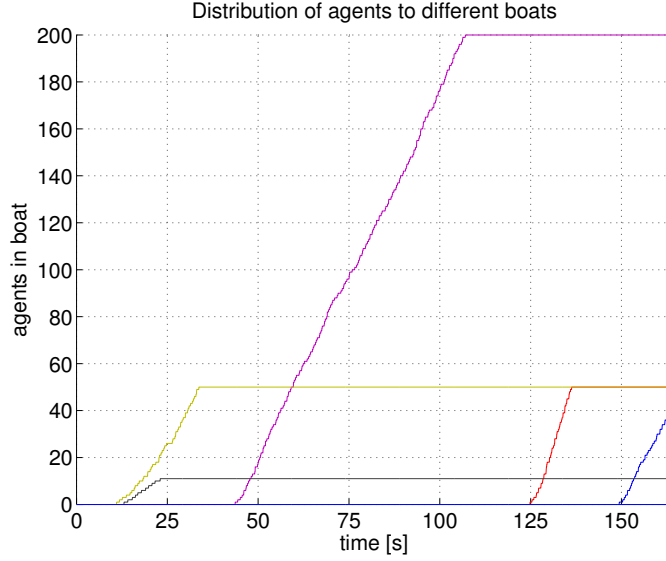


Figure 3: Plot of the filling of the rescue Boats on a ship with 350 Agents, $v_0 = 1.5 \frac{m}{s}$. Each line represents a different rescue boat.

agents stuck in Walls at $v_0 = 1.4m/s$, so we looked at the results, and set the finish time to the value, when the last agent before them left the ship. This introduces an error, which only gets bigger with higher values of v_0 .

To circumvent this, a simulation with a smaller timestep was run, but even with the timestep being a millisecond, we still got agents stuck in walls. An example is shown in Figure 6. This is a part during the simulation with $v_0 = 3.8m/s$ and a timestep of a millisecond.

A timestep of a millisecond means, that an agent, moving at five meters per second, goes a distance of 5 millimeters per timestep. Since the agents get pushed into walls at these rather precise steps of movement, an even smaller timestep will likely yield similar results.

Since we cannot answer our question, whether a high panic level (which means a high value for v_0) with few people or a low panic level with many people is faster, we try to find reasons for this outcome.

A timestep too big can be excluded from the reasons, as we have seen before. We chose the parameters to be the same as in [3] on page 488. They have been found suitable to simulate people leaving a room. The difference to our case is, that we have more walls nearby and because of that, we get a situation where more pressure is applied from more sides. Since that is a remarkable difference to the situation of a single room with one door, it is plausible, that the parameters are not suited for

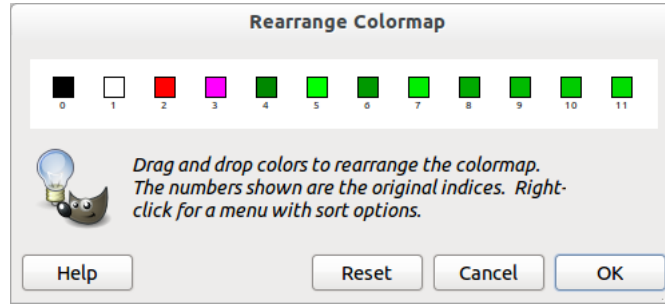


Figure 4: Platzhalter

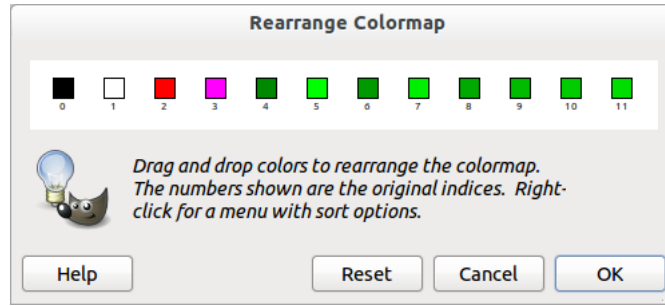


Figure 5: Platzhalter

high densities of people in narrow places.

Of course, this could also be interpreted as people being hurt, but since we did not take this into account for our model, no relevant statements can be made. It should be part of a next simulation to account for the sum of physical forces acting upon an agent. If they are too big, the agent gets hurt. In our case, this would be a realistic assumption, since a ship with limited exit possibilities can make the people feel trapped rather fast and is so more likely to trigger panics than other places.

6.3 Closed Exits

In the simulation, two ways were defined to react to a full boat. First, there was an instantaneous update on the whole ship. This is a model for an announcement over speakers over the whole ship. The agents adapt their direction immediately, which can be seen on a video of the simulation. This behaviour is realistic in a case when there are not too many people who aren't panicking.

The second way is that the information is spread in a circle around the exit. This shall model a simple communication between agents that takes time. With this simple rule of updating the directions, the behaviour is much more realistic, even

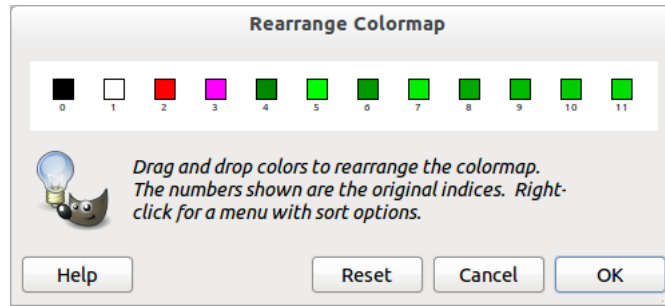


Figure 6: Platzhalter

though the rule doesn't account for the number of people nearby. This means that the information spreads, even if there is nobody. However, an interesting phenomenon during the updates can be observed: While the people near the exit try to get to the next one, they get pushed back by the others who don't know of the change yet.

The flaw of this is that a slow expansion rate leads to more pushing, but if it's too slow, even an agent that escaped can be run against the border of the expansion circle. This is then seen as the agent stopping, even if there is nothing blocking his path. That happens because the agent is faster than the expansion rate of the information. From observations in the simulation, we found $0.5 \frac{m}{s}$ to look mostly like we expected it, even though occasionally there are still some agents slowed down by the expansion rate.

The effect of the pushback on the evacuation time is very small, because of numerous bottlenecks before the actual exit.

TODO: PLOT EINFÜGEN

7 Summary and Outlook

8 References

- [1] <http://www.bbc.co.uk/news/world-europe-16563562>, 9.12.2012
- [2] *Modelling Situations of Evacuation in a Multi-level Building*, Hans Hardmeier, Andrin Jenal, Beat Kng, Felix Thaler, Zurich, April 2012
- [3] *Simulating dynamical features of escape panic*, Dirk Helbing, Illés Farkas, Tamás Vicsek, Nature, 28. September 2000
- [4] <http://www.kreuzfahrtberater.de/deckplan.php?schiff=Costa+Voyager&bf&dpe=2>

- [5] *Pedestrian Dynamics Airplane Evacuation Simulation*, Philipp Heer, Lukas Böhler, Zurich, May 2011
- [6] <http://www.gimp.org/>, 9.12.2012